

Business Context

Renewable energy sources play an increasingly important role in the global energy mix, as the effort to reduce the environmental impact of energy production increases.

Out of all the renewable energy alternatives, wind energy is one of the most developed technologies worldwide. The U.S Department of Energy has put together a guide to achieving operational efficiency using predictive maintenance practices.

Predictive maintenance uses sensor information and analysis methods to measure and predict degradation and future component capability. The idea behind predictive maintenance is that failure patterns are predictable and if component failure can be predicted accurately and the component is replaced before it fails, the costs of operation and maintenance will be much lower.

The sensors fitted across different machines involved in the process of energy generation collect data related to various environmental factors (temperature, humidity, wind speed, etc.) and additional features related to various parts of the wind turbine (gearbox, tower, blades, break, etc.).

Objective

“ReneWind” is a company working on improving the machinery/processes involved in the production of wind energy using machine learning and has collected data of generator failure of wind turbines using sensors. They have shared a ciphered version of the data, as the data collected through sensors is confidential (the type of data collected varies with companies). Data has 40 predictors, 20000 observations in the training set and 5000 in the test set.

The objective is to build various classification models, tune them, and find the best one that will help identify failures so that the generators could be repaired before failing/breaking to reduce the overall maintenance cost. The nature of predictions made by the classification model will translate as follows:

- True positives (TP) are failures correctly predicted by the model. These will result in repairing costs.
- False negatives (FN) are real failures where there is no detection by the model. These will result in replacement costs.
- False positives (FP) are detections where there is no failure. These will result in inspection costs. It is given that the cost of repairing a generator is much less than the cost of replacing it, and the cost of inspection is less than the cost of repair.

“1” in the target variables should be considered as “failure” and “0” represents “No failure”.

Data Description

- The data provided is a transformed version of original data which was collected using sensors.
- Train.csv - To be used for training and tuning of models.
- Test.csv - To be used only for testing the performance of the final best model.
- Both the datasets consist of 40 predictor variables and 1 target variable

```
# Data manipulation and analysis libraries
```

```
import pandas as pd
import numpy as np
```

```
# Data visualization libraries
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Preprocessing and model evaluation libraries
```

```
from sklearn import metrics
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.metrics import (
    f1_score, accuracy_score, recall_score, precision_score,
    confusion_matrix, roc_auc_score
)
from sklearn.model_selection import (
    train_test_split, StratifiedKFold, cross_val_score, RandomizedSearchCV, GridSearchCV
)
```

```
# Sampling libraries
```

```
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
```

```
# Pipeline and transformation libraries
```

```
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
```

```
# Model building libraries
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import (
    AdaBoostClassifier, GradientBoostingClassifier, RandomForestClassifier, BaggingClassifier
)
from xgboost import XGBClassifier
```

```

# Settings to enhance output readability and suppress warnings
pd.set_option("display.max_columns", None)
pd.set_option("display.max_rows", None)
pd.set_option("display.float_format", lambda x: "%.3f" % x) # Suppress scientific notation

# Suppress warnings
import warnings
warnings.filterwarnings("ignore")

train_data = pd.read_csv('./Train.csv')
test_data = pd.read_csv('./Test.csv')

```

Data Overview

```

# View top 5 rows of the data
train_data.head()

```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	\
0	-4.465	-4.679	3.102	0.506	-0.221	-2.033	-2.911	0.051	-1.522	3.762	
1	3.366	3.653	0.910	-1.368	0.332	2.359	0.733	-4.332	0.566	-0.101	
2	-3.832	-5.824	0.634	-2.419	-1.774	1.017	-2.099	-3.173	-2.082	5.393	
3	1.618	1.888	7.046	-1.147	0.083	-1.530	0.207	-2.494	0.345	2.119	
4	-0.111	3.872	-3.758	-2.983	3.793	0.545	0.205	4.849	-1.855	-6.220	

	V11	V12	V13	V14	V15	V16	V17	V18	V19	V20	\
0	-5.715	0.736	0.981	1.418	-3.376	-3.047	0.306	2.914	2.270	4.395	
1	1.914	-0.951	-1.255	-2.707	0.193	-4.769	-2.205	0.908	0.757	-5.834	
2	-0.771	1.107	1.144	0.943	-3.164	-4.248	-4.039	3.689	3.311	1.059	
3	-3.053	0.460	2.705	-0.636	-0.454	-3.174	-3.404	-1.282	1.582	-1.952	
4	1.998	4.724	0.709	-1.989	-2.633	4.184	2.245	3.734	-6.313	-5.380	

	V21	V22	V23	V24	V25	V26	V27	V28	V29	V30	\
0	-2.388	0.646	-1.191	3.133	0.665	-2.511	-0.037	0.726	-3.982	-1.073	
1	-3.065	1.597	-1.757	1.766	-0.267	3.625	1.500	-0.586	0.783	-0.201	
2	-2.143	1.650	-1.661	1.680	-0.451	-4.551	3.739	1.134	-2.034	0.841	
3	-3.517	-1.206	-5.628	-1.818	2.124	5.295	4.748	-2.309	-3.963	-6.029	
4	-0.887	2.062	9.446	4.490	-3.945	4.582	-8.780	-3.383	5.107	6.788	

	V31	V32	V33	V34	V35	V36	V37	V38	V39	V40	\
0	1.667	3.060	-1.690	2.846	2.235	6.667	0.444	-2.369	2.951	-3.480	
1	0.025	-1.795	3.033	-2.468	1.895	-2.298	-1.731	5.909	-0.386	0.616	
2	-1.600	-0.257	0.804	4.086	2.292	5.361	0.352	2.940	3.839	-4.309	
3	4.949	-3.584	-2.577	1.364	0.623	5.550	-1.527	0.139	3.101	-1.277	
4	2.044	8.266	6.629	-10.069	1.223	-3.230	1.687	-2.164	-3.645	6.510	

Target

```

0      0
1      0
2      0
3      0
4      0

# View top 5 rows of the data
test_data.head()

      V1      V2      V3      V4      V5      V6      V7      V8      V9      V10 \
0 -0.613 -3.820  2.202  1.300 -1.185 -4.496 -1.836  4.723  1.206 -0.342
1  0.390 -0.512  0.527 -2.577 -1.017  2.235 -0.441 -4.406 -0.333  1.967
2 -0.875 -0.641  4.084 -1.590  0.526 -1.958 -0.695  1.347 -1.732  0.466
3  0.238  1.459  4.015  2.534  1.197 -3.117 -0.924  0.269  1.322  0.702
4  5.828  2.768 -1.235  2.809 -1.642 -1.407  0.569  0.965  1.918 -2.775

      V11      V12      V13      V14      V15      V16      V17      V18      V19      V20 \
0 -5.123  1.017  4.819  3.269 -2.984  1.387  2.032 -0.512 -1.023  7.339
1  1.797  0.410  0.638 -1.390 -1.883 -5.018 -3.827  2.418  1.762 -3.242
2 -4.928  3.565 -0.449 -0.656 -0.167 -1.630  2.292  2.396  0.601  1.794
3 -5.578 -0.851  2.591  0.767 -2.391 -2.342  0.572 -0.934  0.509  1.211
4 -0.530  1.375 -0.651 -1.679 -0.379 -4.443  3.894 -0.608  2.945  0.367

      V21      V22      V23      V24      V25      V26      V27      V28      V29      V30      V31 \
0 -2.242  0.155  2.054 -2.772  1.851 -1.789 -0.277 -1.255 -3.833 -1.505  1.587
1 -3.193  1.857 -1.708  0.633 -0.588  0.084  3.014 -0.182  0.224  0.865 -1.782
2 -2.120  0.482 -0.841  1.790  1.874  0.364 -0.169 -0.484 -2.119 -2.157  2.907
3 -3.260  0.105 -0.659  1.498  1.100  4.143 -0.248 -1.137 -5.356 -4.546  3.809
4 -5.789  4.598  4.450  3.225  0.397  0.248 -2.362  1.079 -0.473  2.243 -3.591

      V32      V33      V34      V35      V36      V37      V38      V39      V40      Target
0  2.291 -5.411  0.870  0.574  4.157  1.428 -10.511  0.455 -1.448      0
1 -2.475  2.494  0.315  2.059  0.684 -0.485  5.128  1.721 -1.488      0
2 -1.319 -2.997  0.460  0.620  5.632  1.324 -1.752  1.808  1.676      0
3  3.518 -3.074 -0.284  0.955  3.029 -1.367 -3.412  0.906 -2.451      0
4  1.774 -1.502 -2.227  4.777 -6.560 -0.806 -0.276 -3.858 -0.538      0

# Check the dimensions of the data
train_data.shape

(20000, 41)

# Check the dimensions of the data
test_data.shape

(5000, 41)

# Check data types
train_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 41 columns):

```

#	Column	Non-Null	Count	Dtype
0	V1	19982	non-null	float64
1	V2	19982	non-null	float64
2	V3	20000	non-null	float64
3	V4	20000	non-null	float64
4	V5	20000	non-null	float64
5	V6	20000	non-null	float64
6	V7	20000	non-null	float64
7	V8	20000	non-null	float64
8	V9	20000	non-null	float64
9	V10	20000	non-null	float64
10	V11	20000	non-null	float64
11	V12	20000	non-null	float64
12	V13	20000	non-null	float64
13	V14	20000	non-null	float64
14	V15	20000	non-null	float64
15	V16	20000	non-null	float64
16	V17	20000	non-null	float64
17	V18	20000	non-null	float64
18	V19	20000	non-null	float64
19	V20	20000	non-null	float64
20	V21	20000	non-null	float64
21	V22	20000	non-null	float64
22	V23	20000	non-null	float64
23	V24	20000	non-null	float64
24	V25	20000	non-null	float64
25	V26	20000	non-null	float64
26	V27	20000	non-null	float64
27	V28	20000	non-null	float64
28	V29	20000	non-null	float64
29	V30	20000	non-null	float64
30	V31	20000	non-null	float64
31	V32	20000	non-null	float64
32	V33	20000	non-null	float64
33	V34	20000	non-null	float64
34	V35	20000	non-null	float64
35	V36	20000	non-null	float64
36	V37	20000	non-null	float64
37	V38	20000	non-null	float64
38	V39	20000	non-null	float64
39	V40	20000	non-null	float64
40	Target	20000	non-null	int64

```
dtypes: float64(40), int64(1)
memory usage: 6.3 MB
```

```
# Check data types
```

```
test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5000 entries, 0 to 4999
```

```
Data columns (total 41 columns):
```

#	Column	Non-Null Count	Dtype
0	V1	4995 non-null	float64
1	V2	4994 non-null	float64
2	V3	5000 non-null	float64
3	V4	5000 non-null	float64
4	V5	5000 non-null	float64
5	V6	5000 non-null	float64
6	V7	5000 non-null	float64
7	V8	5000 non-null	float64
8	V9	5000 non-null	float64
9	V10	5000 non-null	float64
10	V11	5000 non-null	float64
11	V12	5000 non-null	float64
12	V13	5000 non-null	float64
13	V14	5000 non-null	float64
14	V15	5000 non-null	float64
15	V16	5000 non-null	float64
16	V17	5000 non-null	float64
17	V18	5000 non-null	float64
18	V19	5000 non-null	float64
19	V20	5000 non-null	float64
20	V21	5000 non-null	float64
21	V22	5000 non-null	float64
22	V23	5000 non-null	float64
23	V24	5000 non-null	float64
24	V25	5000 non-null	float64
25	V26	5000 non-null	float64
26	V27	5000 non-null	float64
27	V28	5000 non-null	float64
28	V29	5000 non-null	float64
29	V30	5000 non-null	float64
30	V31	5000 non-null	float64
31	V32	5000 non-null	float64
32	V33	5000 non-null	float64
33	V34	5000 non-null	float64
34	V35	5000 non-null	float64
35	V36	5000 non-null	float64

```

36 V37      5000 non-null   float64
37 V38      5000 non-null   float64
38 V39      5000 non-null   float64
39 V40      5000 non-null   float64
40 Target   5000 non-null   int64
dtypes: float64(40), int64(1)
memory usage: 1.6 MB

# Double check for null values per column
train_data.isnull().sum()

V1          18
V2          18
V3           0
V4           0
V5           0
V6           0
V7           0
V8           0
V9           0
V10          0
V11          0
V12          0
V13          0
V14          0
V15          0
V16          0
V17          0
V18          0
V19          0
V20          0
V21          0
V22          0
V23          0
V24          0
V25          0
V26          0
V27          0
V28          0
V29          0
V30          0
V31          0
V32          0
V33          0
V34          0
V35          0
V36          0

```

```

V37      0
V38      0
V39      0
V40      0
Target    0
dtype: int64

# Checking for null values per column
test_data.isnull().sum()

V1        5
V2        6
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
V29       0
V30       0
V31       0
V32       0
V33       0
V34       0
V35       0
V36       0
V37       0

```



```

V38      0
V39      0
V40      0
Target    0
dtype: int64

# Check data duplicate values
train_data.duplicated().sum()

np.int64(0)

# Check data duplicate values
test_data.duplicated().sum()

np.int64(0)

train_data.describe().T

```

	count	mean	std	min	25%	50%	75%	max
V1	19982.000	-0.272	3.442	-11.876	-2.737	-0.748	1.840	15.493
V2	19982.000	0.440	3.151	-12.320	-1.641	0.472	2.544	13.089
V3	20000.000	2.485	3.389	-10.708	0.207	2.256	4.566	17.091
V4	20000.000	-0.083	3.432	-15.082	-2.348	-0.135	2.131	13.236
V5	20000.000	-0.054	2.105	-8.603	-1.536	-0.102	1.340	8.134
V6	20000.000	-0.995	2.041	-10.227	-2.347	-1.001	0.380	6.976
V7	20000.000	-0.879	1.762	-7.950	-2.031	-0.917	0.224	8.006
V8	20000.000	-0.548	3.296	-15.658	-2.643	-0.389	1.723	11.679
V9	20000.000	-0.017	2.161	-8.596	-1.495	-0.068	1.409	8.138
V10	20000.000	-0.013	2.193	-9.854	-1.411	0.101	1.477	8.108
V11	20000.000	-1.895	3.124	-14.832	-3.922	-1.921	0.119	11.826
V12	20000.000	1.605	2.930	-12.948	-0.397	1.508	3.571	15.081
V13	20000.000	1.580	2.875	-13.228	-0.224	1.637	3.460	15.420
V14	20000.000	-0.951	1.790	-7.739	-2.171	-0.957	0.271	5.671
V15	20000.000	-2.415	3.355	-16.417	-4.415	-2.383	-0.359	12.246
V16	20000.000	-2.925	4.222	-20.374	-5.634	-2.683	-0.095	13.583
V17	20000.000	-0.134	3.345	-14.091	-2.216	-0.015	2.069	16.756
V18	20000.000	1.189	2.592	-11.644	-0.404	0.883	2.572	13.180
V19	20000.000	1.182	3.397	-13.492	-1.050	1.279	3.493	13.238
V20	20000.000	0.024	3.669	-13.923	-2.433	0.033	2.512	16.052
V21	20000.000	-3.611	3.568	-17.956	-5.930	-3.533	-1.266	13.840
V22	20000.000	0.952	1.652	-10.122	-0.118	0.975	2.026	7.410
V23	20000.000	-0.366	4.032	-14.866	-3.099	-0.262	2.452	14.459
V24	20000.000	1.134	3.912	-16.387	-1.468	0.969	3.546	17.163
V25	20000.000	-0.002	2.017	-8.228	-1.365	0.025	1.397	8.223
V26	20000.000	1.874	3.435	-11.834	-0.338	1.951	4.130	16.836
V27	20000.000	-0.612	4.369	-14.905	-3.652	-0.885	2.189	17.560
V28	20000.000	-0.883	1.918	-9.269	-2.171	-0.891	0.376	6.528
V29	20000.000	-0.986	2.684	-12.579	-2.787	-1.176	0.630	10.722
V30	20000.000	-0.016	3.005	-14.796	-1.867	0.184	2.036	12.506

V31	20000.000	0.487	3.461	-13.723	-1.818	0.490	2.731	17.255
V32	20000.000	0.304	5.500	-19.877	-3.420	0.052	3.762	23.633
V33	20000.000	0.050	3.575	-16.898	-2.243	-0.066	2.255	16.692
V34	20000.000	-0.463	3.184	-17.985	-2.137	-0.255	1.437	14.358
V35	20000.000	2.230	2.937	-15.350	0.336	2.099	4.064	15.291
V36	20000.000	1.515	3.801	-14.833	-0.944	1.567	3.984	19.330
V37	20000.000	0.011	1.788	-5.478	-1.256	-0.128	1.176	7.467
V38	20000.000	-0.344	3.948	-17.375	-2.988	-0.317	2.279	15.290
V39	20000.000	0.891	1.753	-6.439	-0.272	0.919	2.058	7.760
V40	20000.000	-0.876	3.012	-11.024	-2.940	-0.921	1.120	10.654
Target	20000.000	0.056	0.229	0.000	0.000	0.000	0.000	1.000

test_data.describe().T

	count	mean	std	min	25%	50%	75%	max
V1	4995.000	-0.278	3.466	-12.382	-2.744	-0.765	1.831	13.504
V2	4994.000	0.398	3.140	-10.716	-1.649	0.427	2.444	14.079
V3	5000.000	2.552	3.327	-9.238	0.315	2.260	4.587	15.315
V4	5000.000	-0.049	3.414	-14.682	-2.293	-0.146	2.166	12.140
V5	5000.000	-0.080	2.111	-7.712	-1.615	-0.132	1.341	7.673
V6	5000.000	-1.042	2.005	-8.924	-2.369	-1.049	0.308	5.068
V7	5000.000	-0.908	1.769	-8.124	-2.054	-0.940	0.212	7.616
V8	5000.000	-0.575	3.332	-12.253	-2.642	-0.358	1.713	10.415
V9	5000.000	0.030	2.174	-6.785	-1.456	-0.080	1.450	8.851
V10	5000.000	0.019	2.145	-8.171	-1.353	0.166	1.511	6.599
V11	5000.000	-2.009	3.112	-13.152	-4.050	-2.043	0.044	9.956
V12	5000.000	1.576	2.907	-8.164	-0.450	1.488	3.563	12.984
V13	5000.000	1.622	2.883	-11.548	-0.126	1.719	3.465	12.620
V14	5000.000	-0.921	1.803	-7.814	-2.111	-0.896	0.272	5.734
V15	5000.000	-2.452	3.387	-15.286	-4.479	-2.417	-0.433	11.673
V16	5000.000	-3.019	4.264	-20.986	-5.648	-2.774	-0.178	13.976
V17	5000.000	-0.104	3.337	-13.418	-2.228	0.047	2.112	19.777
V18	5000.000	1.196	2.586	-12.214	-0.409	0.881	2.604	13.642
V19	5000.000	1.210	3.385	-14.170	-1.026	1.296	3.526	12.428
V20	5000.000	0.138	3.657	-13.720	-2.325	0.193	2.540	13.871
V21	5000.000	-3.664	3.578	-16.341	-5.944	-3.663	-1.330	11.047
V22	5000.000	0.962	1.640	-6.740	-0.048	0.986	2.029	7.505
V23	5000.000	-0.422	4.057	-14.422	-3.163	-0.279	2.426	13.181
V24	5000.000	1.089	3.968	-12.316	-1.623	0.913	3.537	17.806
V25	5000.000	0.061	2.010	-6.770	-1.298	0.077	1.428	6.557
V26	5000.000	1.847	3.400	-11.414	-0.242	1.917	4.156	17.528
V27	5000.000	-0.552	4.403	-13.177	-3.663	-0.872	2.247	17.290
V28	5000.000	-0.868	1.926	-7.933	-2.160	-0.931	0.421	7.416
V29	5000.000	-1.096	2.655	-9.988	-2.861	-1.341	0.522	14.039
V30	5000.000	-0.119	3.023	-12.438	-1.997	0.112	1.946	10.315
V31	5000.000	0.469	3.446	-11.263	-1.822	0.486	2.779	12.559
V32	5000.000	0.233	5.586	-17.244	-3.556	-0.077	3.752	26.539

```

V33      5000.000 -0.080 3.539 -14.904 -2.348 -0.160  2.099 13.324
V34      5000.000 -0.393 3.166 -14.700 -2.010 -0.172  1.465 12.146
V35      5000.000  2.211 2.948 -12.261  0.322  2.112  4.032 13.489
V36      5000.000  1.595 3.775 -12.736 -0.866  1.703  4.104 17.116
V37      5000.000  0.023 1.785  -5.079 -1.241 -0.110  1.238  6.810
V38      5000.000 -0.406 3.969 -15.335 -2.984 -0.381  2.288 13.065
V39      5000.000  0.939 1.717  -5.451 -0.208  0.959  2.131  7.182
V40      5000.000 -0.932 2.978 -10.076 -2.987 -1.003  1.080  8.698
Target  5000.000  0.056 0.231   0.000  0.000  0.000  0.000  1.000

```

```

# Check for class distribution of the target variable by looking at the count of observation
class_distribution = train_data['Target'].value_counts(normalize=True) * 100
print(class_distribution)

```

```

Target
0    94.450
1     5.550
Name: proportion, dtype: float64

```

```

# Check for class distribution of the target variable by looking at the count of observation
class_distribution = test_data['Target'].value_counts(normalize=True) * 100
print(class_distribution)

```

```

Target
0    94.360
1     5.640
Name: proportion, dtype: float64

```

Observations:

Dataset Size:

- The training set contains 20,000 observations, each with 41 columns (40 predictors and 1 target variable).
- The test set includes 5,000 observations, also with 41 columns.

Data Types:

- Both datasets primarily consist of numerical features (float64), with the target variable being an integer (int64).

Missing Values:

- In the training set, two variables (V1 and V2) have 18 missing values each.
- In the test set, V1 has 5 missing values, and V2 has 6 missing values.
- No missing values are present in the target variable of both datasets.

Data Integrity:

- There are no duplicate rows in either the training or the test dataset, indicating good data integrity.

Class Imbalance: Bot training and testing target variables have imbalanced distribution of the target variables classes where:

- About 94% of the observations belong to the negative class (0), which represents "Non-failure" cases.
- About 5% of the observations belong to the positive class (1), representing "Failure" cases.

Exploratory Data analysis (EDA)

Univariate Analysis

```
def hist_and_boxplot(data, variable, figsize=(12, 4), kde=False, bins=None):
    """
    Creates a plot with both a histogram and boxplot for a specified numerical variable.

    Args:
    - data: The DataFrame containing the data.
    - variable: Column name of the numerical variable (feature) to be plotted.
    - figsize: A tuple representing the size of the figure.
    - density_curve: A boolean indicating whether to overlay a density curve on the histogram.
    - bins: An integer representing the number of bins for the histogram, or None for automatic.

    Returns:
    None
    """
    # Set up the matplotlib figure with two rows and one column
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=figsize, sharex=True, gridspec_kw={'height_ratios': [1, 1]})

    # Plot the boxplot on the first row
    sns.boxplot(x=variable, data=data, ax=ax1, showmeans=True, color="lightblue")
    ax1.set(xlabel='', title=f'Boxplot and Distribution of {variable}')

    # Plot the histogram on the second row
    if bins:
        sns.histplot(data[variable], kde=kde, bins=bins, ax=ax2, color="lightblue")
    else:
        sns.histplot(data[variable], kde=kde, ax=ax2, color="lightblue")

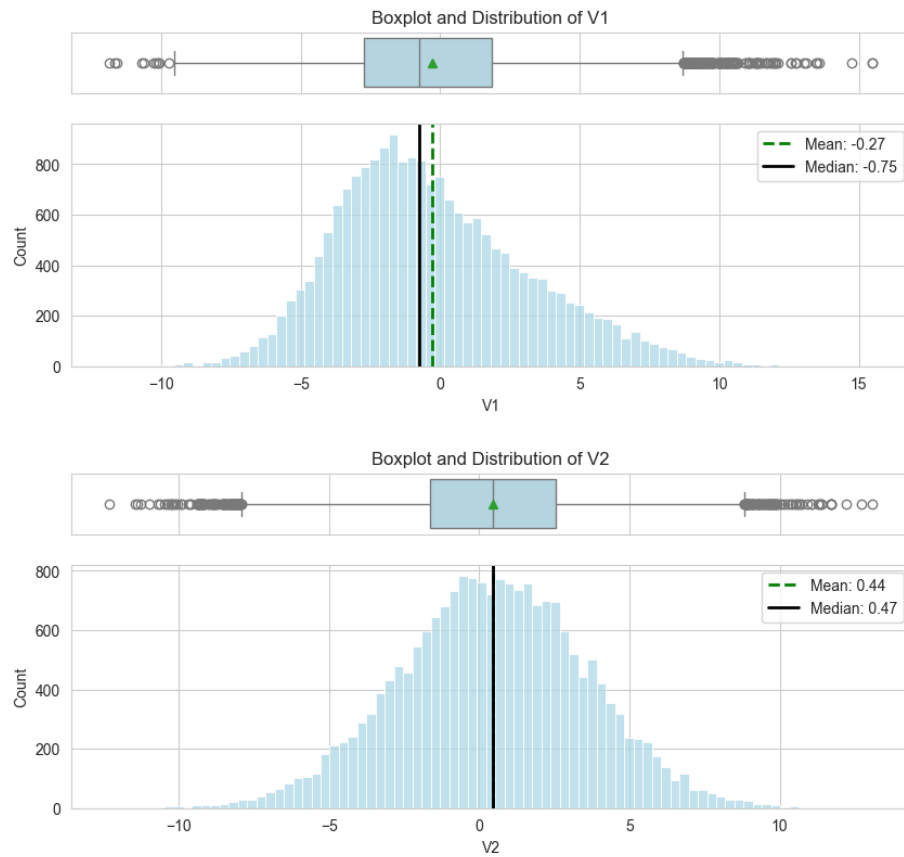
    # Draw lines for mean and median
    mean_val = data[variable].mean()
    median_val = data[variable].median()
    ax2.axvline(mean_val, color='green', linestyle='--', linewidth=2, label=f'Mean: {mean_val}')
    ax2.axvline(median_val, color='black', linestyle='-', linewidth=2, label=f'Median: {median_val}')


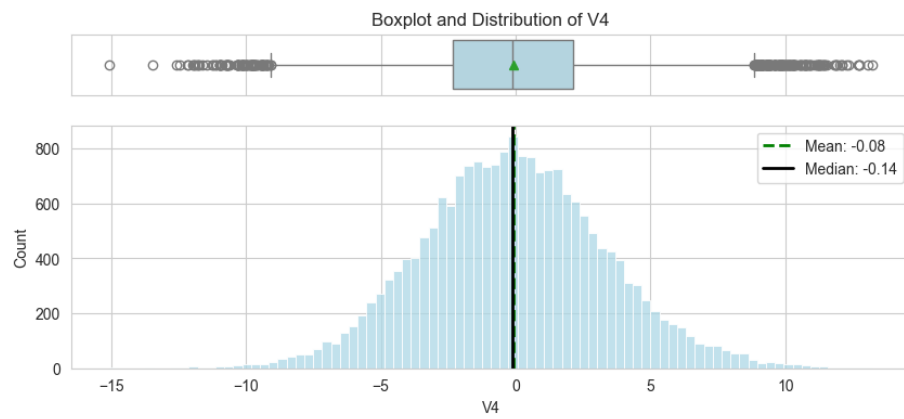

    # Add legend to the histogram
```




```
ax2.legend()
```

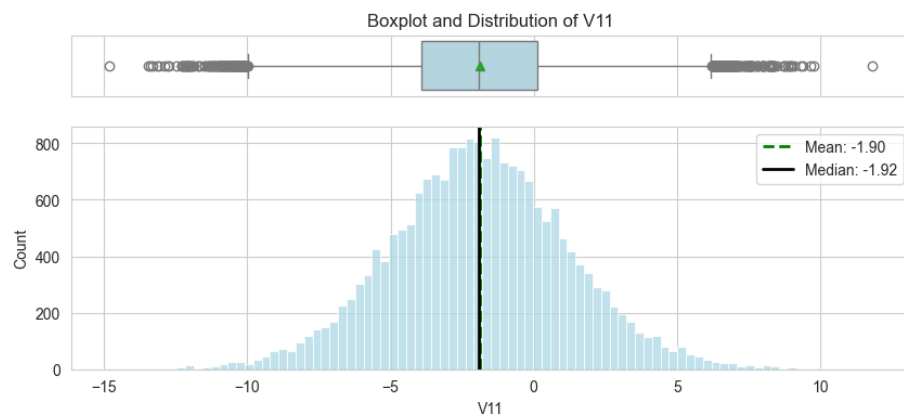


```
plt.show()
```




```
# Looping through the variables in the train dataset and creating a histogram and boxplot for  
for variable in train_data.columns:  
    hist_and_boxplot(train_data, variable, figsize=(10, 4), kde=False, bins=None)
```































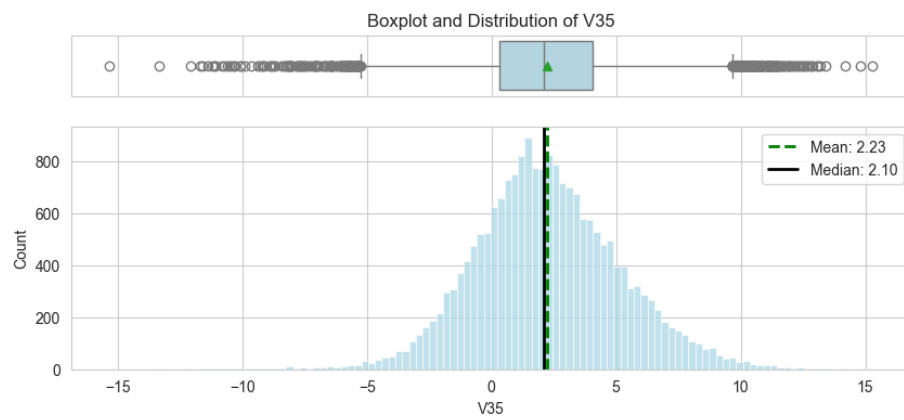









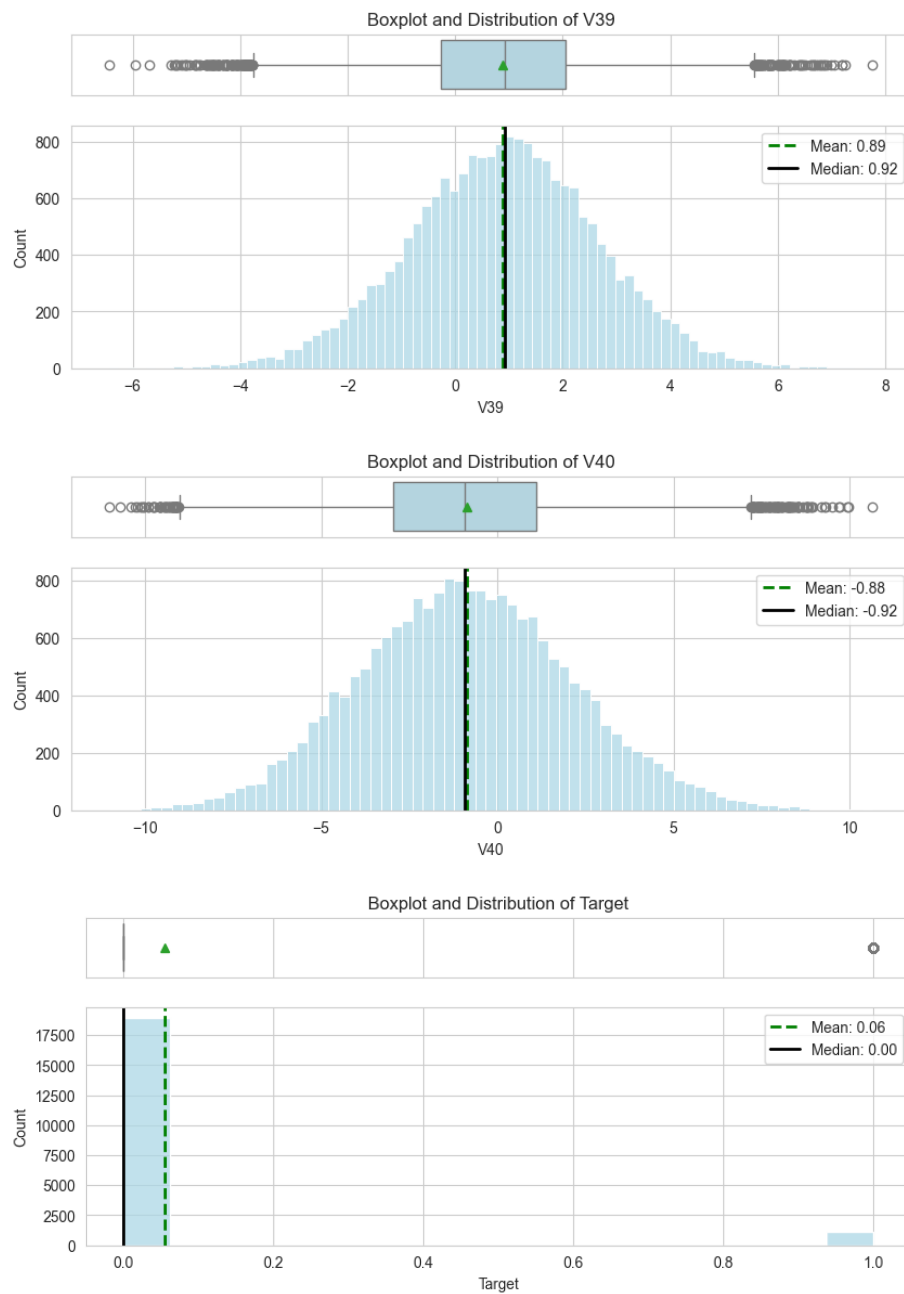












Observations:

Independent variables:

- The histograms indicate that the data for these variables is approximately

normally distributed with mean and median fairly close for most, further suggesting a normal distribution.

- The boxplots show dots outside of the whiskers, which represent outliers. These are observations that fall significantly higher or lower than the majority of the data.


The target variable:

- Highly imbalanced, with a much larger count of the 0 class compared to the 1 class consistnat with the numerical class distribution observation noted earlier.

Bivariate Analysis

Heatmap of the correlation coefficients between numerical variables in the dataset
`cols_list = train_data.select_dtypes(include=np.number).columns.tolist()`

```
plt.figure(figsize=(18, 12))
sns.heatmap(
    train_data[cols_list].corr(), annot=True, vmin=-1, vmax=1, fmt=".1f", cmap="coolwarm"
)
plt.show()
```



Observations:

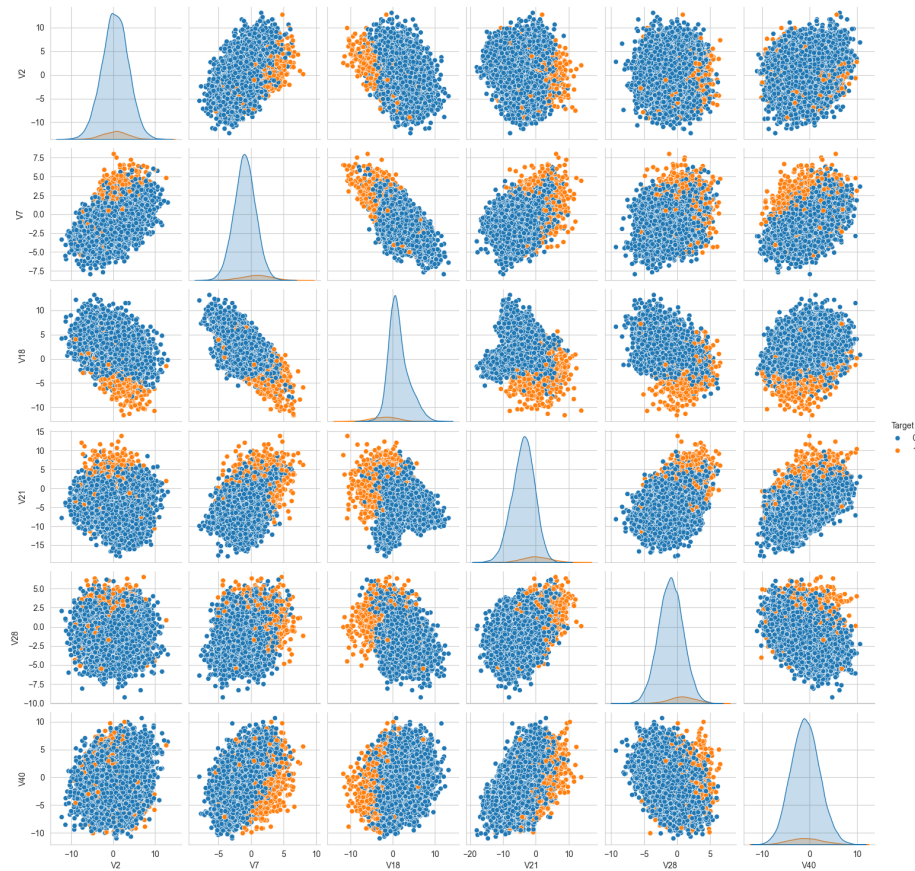
- Highly Correlated Features: lots of them (both positive and negative). These

may represent redundancy within the dataset which means that one feature can be predicted from the other with a high degree of accuracy. In other words, these features carry similar information or signals about the data.

- Multicollinearity in liner models
- Overfitting the ti the training dataset due to redundant features
- Redundant features increase computational complexity
- Consider dimensionality reduction techniques like PCA
- Relationship with Target: features that have a stronger correlation with the target variable (V18, V21). These features may be particularly important for predicting the target and should be examined more closely in subsequent analyses.

Selected a subset of features for the pair plot for better performance and clarity
selected_features = ['V2', 'V7', 'V18', 'V21', 'V28', 'V40', 'Target']

Create the pair plot using the selected features
sns.pairplot(train_data[selected_features], hue='Target', diag_kind='kde')
plt.show()



```
train_df = train_data.copy()
test_df = test_data.copy()
```

Data Pre-Processing

```
# Split train data into X and y to separate the features from the target
X = train_df.drop(["Target"], axis=1)
y = train_df["Target"]

# Split the X training set into train and validate
X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.25, random_state=42)

# Check the number of rows and columns in the X_train data
X_train.shape

(15000, 40)

# Check the number of rows and columns in the X_val data
X_val.shape

(5000, 40)

# Split test data into X and y to separate the features from the target
X_test = test_df.drop(["Target"], axis=1)
y_test = test_df["Target"]

# Check the number of rows and columns in the X_test data
X_test.shape

(5000, 40)
```

- X_train and y_train are now the subsets for training the model.
- X_val and y_val are for validating the model during the hyperparameter tuning and model selection process.
- X_test and y_test (from the original test dataset) are used strictly for the final evaluation of the model to assess its performance on unseen data.

Missing Value Imputation

```
# Create an instance of the imputer
imputer = SimpleImputer(strategy="median")

# Fit on the training data and transform it
X_train = pd.DataFrame(imputer.fit_transform(X_train), columns=X_train.columns)

# Transform the validation data based on the fit from the training data
X_val = pd.DataFrame(imputer.transform(X_val), columns=X_val.columns)
```

```

# Transform the test data based on the fit from the training data
X_test = pd.DataFrame(imputer.transform(X_test), columns=X_test.columns)

# Check that it worked
print(X_train.isna().sum())
print("-"*15)
print(X_val.isna().sum())
print("-"*15)
print(X_test.isna().sum())
print("-"*15)

```

V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0
V24	0
V25	0
V26	0
V27	0
V28	0
V29	0
V30	0
V31	0
V32	0
V33	0
V34	0
V35	0
V36	0

```
V37    0
V38    0
V39    0
V40    0
dtype: int64
-----
V1      0
V2      0
V3      0
V4      0
V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
V29     0
V30     0
V31     0
V32     0
V33     0
V34     0
V35     0
V36     0
V37     0
V38     0
V39     0
V40     0
```

```
dtype: int64
-----
V1      0
V2      0
V3      0
V4      0
V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
V29     0
V30     0
V31     0
V32     0
V33     0
V34     0
V35     0
V36     0
V37     0
V38     0
V39     0
V40     0
dtype: int64
-----
```


Scale/Normalize Features

```
# Create a scaler instance
scaler = StandardScaler()

# Fit on the training data and transform it
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns=X_train.columns)

# Transform the validation and test data
X_val = pd.DataFrame(scaler.transform(X_val), columns=X_val.columns)

X_test = pd.DataFrame(scaler.transform(X_test), columns=X_test.columns)
```

Model Building

```
# defining a function to compute different metrics to check performance of a classification
def model_performance_classification_sklearn(model, predictors, target):
    """
    Function to compute different metrics to check classification model performance

    model: classifier
    predictors: independent variables
    target: dependent variable
    """
    TP = confusion_matrix(target, model.predict(predictors))[1,1]
    FP = confusion_matrix(target, model.predict(predictors))[0,1]
    FN = confusion_matrix(target, model.predict(predictors))[1,0]

    # predicting using the independent variables
    pred = model.predict(predictors)

    acc = accuracy_score(target, pred) # to compute Accuracy
    recall = recall_score(target, pred) # to compute Recall
    precision = precision_score(target, pred) # to compute Precision
    f1 = f1_score(target, pred) # to compute F1-score

    # creating a dataframe of metrics
    df_perf = pd.DataFrame(
        {
            "Accuracy": acc,
            "Recall": recall,
            "Precision": precision,
            "F1": f1
        },
        index=[0],
```

```
)
```

```
return df_perf
```

Defining scorer for cross-validation and hyperparameter tuning

- The goal here is to reduce FN and maximize "Recall"
- So, use Recall as a scorer

```
# Defining the scorer based on recall
```

```
scorer = metrics.make_scorer(metrics.recall_score)
```

Model Building on Original Data

```
# Models to evaluate
```

```
models = []
```

```
models.append(("DecisionTree", DecisionTreeClassifier(random_state=1)))
```

```
models.append(("RandomForest", RandomForestClassifier(random_state=1)))
```

```
models.append(("AdaBoost", AdaBoostClassifier(random_state=1)))
```

```
models.append(("GradientBoosting", GradientBoostingClassifier(random_state=1)))
```

```
models.append(("BaggingClassifier", BaggingClassifier(random_state=1)))
```

```
models.append(("LogisticRegression", LogisticRegression(random_state=1, max_iter=10000)))
```

```
results_original = [] # To store cross-validation results
```

```
names_original = [] # To store model names
```

```
# Cross-validation across all models for Original Data
```

```
print("\nCross-Validation on Original Data:\n")
```

```
# StratifiedKFold setup
```

```
kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)
```

```
# Cross-validation across all models
```

```
for name, model in models:
```

```
    cv_result = cross_val_score(model, X_train, y_train, scoring=scorer, cv=kfold)
```

```
    results_original.append(cv_result)
```

```
    names_original.append(name)
```

```
    print(f"{name}: Mean Recall Score = {cv_result.mean()}")
```

```
print("\nValidation Performance on Original Data:\n")
```

```
# Fit models on the original training set and evaluate on the original validation set
```

```
for name, model in models:
```

```
    model.fit(X_train, y_train) # Use the training data
```

```

scores = recall_score(y_val, model.predict(X_val)) # Evaluate against the validation set
print(f"{name}: Validation Recall Score = {scores}")

```

Cross-Validation on Original Data:

```

DecisionTree: Mean Recall Score = 0.7297619047619047
RandomForest: Mean Recall Score = 0.7261904761904762
AdaBoost: Mean Recall Score = 0.536904761904762
GradientBoosting: Mean Recall Score = 0.7142857142857142
BaggingClassifier: Mean Recall Score = 0.705952380952381
LogisticRegression: Mean Recall Score = 0.48809523809523814

```

Validation Performance on Original Data:

```

DecisionTree: Validation Recall Score = 0.7111111111111111
RandomForest: Validation Recall Score = 0.6962962962962963
AdaBoost: Validation Recall Score = 0.5333333333333333
GradientBoosting: Validation Recall Score = 0.6888888888888889
BaggingClassifier: Validation Recall Score = 0.7222222222222222
LogisticRegression: Validation Recall Score = 0.48148148148148145

```

```

# Plotting boxplots for CV scores of all models defined above

```

```

fig = plt.figure(figsize=(10, 7))

```

```

fig.suptitle("Algorithm Comparison of Models Trained on Original Data")
ax = fig.add_subplot(111)

```

```

plt.boxplot(results_original)
ax.set_xticklabels(names_original)


```

```

plt.show()

```

Algorithm Comparison of Models Trained on Original Data



Recall is highest for Random Forest followed by Decision Tree and Gradient Boosting.

Model Building with Oversampled Data

```
# Fit Synthetic Minority Over Sampling Technique (SMOTE) on train data, which uses kNN to g
sm = SMOTE(sampling_strategy=1, k_neighbors=5, random_state=1)
X_train_over, y_train_over = sm.fit_resample(X_train, y_train)

print("Before OverSampling, count of label '1': {}".format(sum(y_train == 1)))
print("Before OverSampling, count of label '0': {} \n".format(sum(y_train == 0)))

print("After OverSampling, count of label '1': {}".format(sum(y_train_over == 1)))
print("After OverSampling, count of label '0': {} \n".format(sum(y_train_over == 0)))

print("After OverSampling, the shape of train_X: {}".format(X_train_over.shape))
print("After OverSampling, the shape of train_y: {} \n".format(y_train_over.shape))

Before OverSampling, count of label '1': 840
Before OverSampling, count of label '0': 14160

After OverSampling, count of label '1': 14160
After OverSampling, count of label '0': 14160
```

After OverSampling, the shape of train_X: (28320, 40)
After OverSampling, the shape of train_y: (28320,)

Observations:

SMOTE has successfully balanced your dataset by generating synthetic examples of the minority class (1) using k-nearest neighbors.

Before applying SMOTE: the training dataset was significantly imbalanced with only 833 instances of the minority class (1) compared to 14,167 instances of the majority class (0). This imbalance leads to models that are biased towards the majority class and ignoring the minority class which is of greater interest to us.

After applying SMOTE: both classes now have an equal count of 14,167 instances, doubling the size of the training dataset to 28,334 instances. This balanced dataset is expected to improve model performance on the minority class by providing it with more examples to learn from, thus helping the model to generalize better and be less biased towards the majority class.

```
# Models to evaluate on Oversampled Data
models = []
models.append(("DecisionTree", DecisionTreeClassifier(random_state=1)))
models.append(("RandomForest", RandomForestClassifier(random_state=1)))
models.append(("AdaBoost", AdaBoostClassifier(random_state=1)))
models.append(("GradientBoosting", GradientBoostingClassifier(random_state=1)))
models.append(("BaggingClassifier", BaggingClassifier(random_state=1)))
models.append(("LogisticRegression", LogisticRegression(random_state=1, max_iter=10000)))

# To store cross-validation results
results_oversampled = []
# To store model names
names_oversampled= []

# Cross-validation across all models for Oversampled Data
print("\nCross-Validation on Oversampled Data:\n")

# StratifiedKFold setup
kfold_oversampled = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)

# Cross-validation across all models
for name, model in models:
    cv_result_oversampled = cross_val_score(model, X_train_over, y_train_over, scoring=score
    results_oversampled.append(cv_result_oversampled)
    names_oversampled.append(name)
    print(f"{name}: Mean Recall Score = {cv_result_oversampled.mean()}")
```

```

print("\nValidation Performance on Oversampled Data:\n")

# Fit models on the oversampled training set and evaluate on the original validation set
for name, model in models:
    model.fit(X_train_over, y_train_over) # Use the oversampled training data
    scores_oversampled = recall_score(y_val, model.predict(X_val)) # Evaluate against the
    print(f"{name}: Validation Recall Score = {scores_oversampled}")

```

Cross-Validation on Oversampled Data:

```

DecisionTree: Mean Recall Score = 0.9719632768361581
RandomForest: Mean Recall Score = 0.984039548022599
AdaBoost: Mean Recall Score = 0.8841101694915254
GradientBoosting: Mean Recall Score = 0.9216807909604519
BaggingClassifier: Mean Recall Score = 0.9750706214689266
LogisticRegression: Mean Recall Score = 0.8748587570621469

```

Validation Performance on Oversampled Data:

```

DecisionTree: Validation Recall Score = 0.7814814814814814
RandomForest: Validation Recall Score = 0.8296296296296296
AdaBoost: Validation Recall Score = 0.8444444444444444
GradientBoosting: Validation Recall Score = 0.8814814814814815
BaggingClassifier: Validation Recall Score = 0.8111111111111111
LogisticRegression: Validation Recall Score = 0.8518518518518519

```

Observations:

As expected, the balancing of the data with SMOTE improved model performance accross the board. However, it also introduced overfitting issues (ratio of Cross-Validation to Validation scores). Except for AdaBoost and Logistic Regression.

```

# Plotting boxplots for CV scores of all models evaluated on oversampled data
fig = plt.figure(figsize=(10, 7))

```

```

fig.suptitle("Algorithm Comparison of Models Trained on Oversampled Data")
ax = fig.add_subplot(111)

```

```

plt.boxplot(results_oversampled)
ax.set_xticklabels(names_oversampled)


```

```

plt.show()

```

Algorithm Comparison of Models Trained on Oversampled Data



Observations:

Same trend, jsut more exaggerated.

Model Building with Undersampled Data

```
# Random undersampler for under sampling the data
rus = RandomUnderSampler(random_state=1, sampling_strategy=1)
X_train_un, y_train_un = rus.fit_resample(X_train, y_train)

print("Before Under Sampling, count of label '1': {}".format(sum(y_train== 1)))
print("Before Under Sampling, count of label '0': {} \n".format(sum(y_train == 0)))

print("After Under Sampling, count of label '1': {}".format(sum(y_train_un == 1)))
print("After Under Sampling, count of label '0': {} \n".format(sum(y_train_un == 0)))

print("After Under Sampling, the shape of train_X: {}".format(X_train_un.shape))
print("After Under Sampling, the shape of train_y: {} \n".format(y_train_un.shape))

Before Under Sampling, count of label '1': 840
Before Under Sampling, count of label '0': 14160

After Under Sampling, count of label '1': 840
After Under Sampling, count of label '0': 840
```

After Under Sampling, the shape of train_X: (1680, 40)

After Under Sampling, the shape of train_y: (1680,)

```
# Models to evaluate on Undersampled Data
models_undersampled = []
models_undersampled.append(("DecisionTree", DecisionTreeClassifier(random_state=1)))
models_undersampled.append(("RandomForest", RandomForestClassifier(random_state=1)))
models_undersampled.append(("AdaBoost", AdaBoostClassifier(random_state=1)))
models_undersampled.append(("GradientBoosting", GradientBoostingClassifier(random_state=1)))
models_undersampled.append(("BaggingClassifier", BaggingClassifier(random_state=1)))
models_undersampled.append(("LogisticRegression", LogisticRegression(random_state=1, max_iter=10000)))

# To store cross-validation results for undersampled data
results_undersampled = []
# To store model names for undersampled data
names_undersampled = []

# Cross-validation across all models for Undersampled Data
print("\nCross-Validation on Undersampled Data:\n")

# StratifiedKFold setup for undersampled data
kfold_undersampled = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)

# Cross-validation across all models
for name, model in models_undersampled:
    cv_result_undersampled = cross_val_score(model, X_train_un, y_train_un, scoring=scorer,
                                              cv=kfold_undersampled)
    results_undersampled.append(cv_result_undersampled)
    names_undersampled.append(name)
    print(f"{name}: Mean Recall Score = {cv_result_undersampled.mean()}")

print("\nValidation Performance on Undersampled Data:\n")

# Fit models on the undersampled training set and evaluate on the original validation set
for name, model in models_undersampled:
    model.fit(X_train_un, y_train_un) # Use the undersampled training data
    scores_undersampled = recall_score(y_val, model.predict(X_val)) # Evaluate against the original validation set
    print(f"{name}: Validation Recall Score = {scores_undersampled}")
```

Cross-Validation on Undersampled Data:

DecisionTree: Mean Recall Score = 0.8678571428571427

RandomForest: Mean Recall Score = 0.8988095238095237

AdaBoost: Mean Recall Score = 0.8607142857142858

GradientBoosting: Mean Recall Score = 0.8952380952380953


```
BaggingClassifier: Mean Recall Score = 0.8738095238095237
LogisticRegression: Mean Recall Score = 0.8547619047619047
```

Validation Performance on Undersampled Data:


```
DecisionTree: Validation Recall Score = 0.837037037037037
RandomForest: Validation Recall Score = 0.8777777777777778
AdaBoost: Validation Recall Score = 0.8555555555555555
GradientBoosting: Validation Recall Score = 0.8888888888888888
BaggingClassifier: Validation Recall Score = 0.8481481481481481
LogisticRegression: Validation Recall Score = 0.8555555555555555
```

```
# Plotting boxplots for CV scores of all models evaluated on undersampled data
fig = plt.figure(figsize=(10, 7))
```

```
fig.suptitle("Algorithm Comparison of Models Trained on Undersampled Data")
ax = fig.add_subplot(111)
```

```
plt.boxplot(results_undersampled)
ax.set_xticklabels(names_undersampled)
```

```
plt.show()
```



Observations:

Fast. Random Forest still highest performace. Less overfitting accross the board.

Hyperparameter Tuning

Model Selection for Hyperparameter Tuning

Top 3 models providing the highest Recall scores on both cross-validation and validation datasets :

1. **RandomForest**
2. **GradientBoosting**
3. **AdaBoost**

Hyperparameter Tuning of Random Forest on Original Data

```
# Defining the Random Forest model for original data
rf_model_orig = RandomForestClassifier(random_state=1)

# Parameter grid for Random Forest
param_grid_rf_orig = {
    "n_estimators": [200, 250], # Reduced for simplicity
    "min_samples_leaf": [1, 2], # Simplified range
    "max_features": ['sqrt'], # Simplified choice
    "max_samples": [0.5] # Simplified choice
}

# Setting up RandomizedSearchCV for the original data with reduced n_iter and cv
randomized_rf_orig = RandomizedSearchCV(estimator=rf_model_orig, param_distributions=param_grid_rf_orig,
                                         scoring= scorer, n_iter=10, cv=3, random_state=1)

# Fitting RandomizedSearchCV on the original data
randomized_rf_orig.fit(X_train, y_train)

# Extracting best parameters and creating a new Random Forest Classifier
best_params_rf_orig = randomized_rf_orig.best_params_
tuned_rf_orig = RandomForestClassifier(**best_params_rf_orig, random_state=1)

# Fitting the tuned model and evaluating
tuned_rf_orig.fit(X_train, y_train)

RandomForestClassifier(max_samples=0.5, n_estimators=200, random_state=1)

# Evaluating the tuned model on original data
rf_train_perf_orig = model_performance_classification_sklern(tuned_rf_orig, X_train, y_train)
print("Performance on Train Set:")
rf_train_perf_orig
```

Performance on Train Set:

	Accuracy	Recall	Precision	F1
0	0.993	0.879	0.999	0.935

```
rf_val_perf_orig = model_performance_classification_sklearn(tuned_rf_orig, X_val, y_val)
print("Performance on Validation Set:")
rf_val_perf_orig
```

Performance on Validation Set:

	Accuracy	Recall	Precision	F1
0	0.982	0.670	0.989	0.799

Hyperparameter Tuning of Random Forest on Oversampled Data

```
# Defining the Random Forest model for original data
```

```
rf_model_over = RandomForestClassifier(random_state=1)
```

```
# Parameter grid for Random Forest
```

```
param_grid_rf_orig = {
    "n_estimators": [200, 250], # Reduced for simplicity
    "min_samples_leaf": [1, 2], # Simplified range
    "max_features": ['sqrt'], # Simplified choice
    "max_samples": [0.5] # Simplified choice
}
```

```
# Setting up RandomizedSearchCV for the original data
```

```
randomized_rf_over = RandomizedSearchCV(estimator=rf_model_orig, param_distributions=param_g
                                         scoring=scorer, n_iter=50, cv=5, random_state=1)
```

```
# Fitting RandomizedSearchCV on the oversampled data
```

```
randomized_rf_over.fit(X_train_over, y_train_over)
```

```
# Extracting best parameters for oversampled data
```

```
best_params_rf_over = randomized_rf_over.best_params_
```

```
# Creating a new Random Forest Classifier with the best parameters for oversampled data
```

```
tuned_rf_over = RandomForestClassifier(**best_params_rf_over, random_state=1)
```

```
# Fitting the tuned model on the oversampled training data
```

```
tuned_rf_over.fit(X_train_over, y_train_over)
```

```
RandomForestClassifier(max_samples=0.5, n_estimators=250, random_state=1)
```

```
# Evaluating the tuned model on oversampled train data
```

```
rf_train_perf_over = model_performance_classification_sklearn(tuned_rf_over, X_train_over, y
rf_train_perf_over
```

	Accuracy	Recall	Precision	F1
0	0.999	0.998	1.000	0.999

Evaluating the tuned model on oversampled validation data

```
rf_val_perf_over = model_performance_classification_sklearn(tuned_rf_over, X_val, y_val)
rf_val_perf_over
```

	Accuracy	Recall	Precision	F1
0	0.988	0.859	0.917	0.887

Hyperparameter Tuning of Random Forest on Undersampled Data

Defining the Random Forest model for original data

```
rf_model_un = RandomForestClassifier(random_state=1)
```

Parameter grid for Random Forest

```
param_grid_rf_orig = {
    "n_estimators": [200, 250], # Reduced for simplicity
    "min_samples_leaf": [1, 2], # Simplified range
    "max_features": ['sqrt'], # Simplified choice
    "max_samples": [0.5] # Simplified choice
}
```

Setting up RandomizedSearchCV for the original data

```
randomized_rf_un = RandomizedSearchCV(estimator=rf_model_orig, param_distributions=param_grid_rf_orig,
                                       scoring= scorer, n_iter=50, cv=5, random_state=1)
```

Fitting RandomizedSearchCV on the undersampled data

```
randomized_rf_un.fit(X_train_un, y_train_un)
```

Extracting best parameters for undersampled data

```
best_params_rf_un = randomized_rf_un.best_params_
print("Best parameters are {} with CV score={}" .format(randomized_rf_un.best_params_, randomized_rf_un.best_score_))
```

Best parameters are {'n_estimators': 250, 'min_samples_leaf': 2, 'max_samples': 0.5, 'max_features': 'sqrt'}

Creating a new Random Forest Classifier with the best parameters for undersampled data

```
tuned_rf_un = RandomForestClassifier(**best_params_rf_un, random_state=1)
```

Fitting the tuned model on the undersampled training data

```
tuned_rf_un.fit(X_train_un, y_train_un)
```

```
RandomForestClassifier(max_samples=0.5, min_samples_leaf=2, n_estimators=250,
                       random_state=1)
```

Evaluating the tuned model on undersampled data

```
rf_train_perf_un = model_performance_classification_sklearn(tuned_rf_un, X_train_un, y_train_un)
rf_train_perf_un
```

	Accuracy	Recall	Precision	F1
0	0.966	0.940	0.991	0.965

```
rf_val_perf_un = model_performance_classification_sklearn(tuned_rf_un, X_val, y_val)
rf_val_perf_un
```

	Accuracy	Recall	Precision	F1
0	0.935	0.878	0.450	0.595

Hyperparameter Tuning of Gradient Boost on Original Data

Defining the Gradient Boosting model

```
gb_model_orig = GradientBoostingClassifier(random_state=1)
```

Parameter grid remains the same as the oversampling case

```
param_grid_orig = {
    "n_estimators": np.arange(100, 150, 25),
    "learning_rate": [0.2, 0.05, 1],
    "subsample": [0.5, 0.7],
    "max_features": [0.5, 0.7]
}
```

Setting up RandomizedSearchCV for the original data

```
randomized_cv_orig = RandomizedSearchCV(estimator=gb_model_orig, param_distributions=param_g
                                         scoring=scorer, n_iter=50, n_jobs=-1, cv=5, random_s
```

Fitting RandomizedSearchCV on the original data

```
randomized_cv_orig.fit(X_train, y_train)
```

Best parameters

```
# print("Best parameters for original data are {} with CV score={}".format(randomized_cv_o
```

Extracting best parameters for the original data

```
best_params_orig = randomized_cv_orig.best_params_
```

Creating a new Gradient Boosting Classifier with the best parameters for original data

```
tuned_gbm_orig = GradientBoostingClassifier(
    max_features=best_params_orig['max_features'],
    random_state=1,
    learning_rate=best_params_orig['learning_rate'],
    n_estimators=best_params_orig['n_estimators'],
    subsample=best_params_orig['subsample']
)
```

Fitting the tuned model on the original training data

```
tuned_gbm_orig.fit(X_train, y_train)
```

```

GradientBoostingClassifier(learning_rate=0.2, max_features=0.7,
                           n_estimators=np.int64(125), random_state=1,
                           subsample=0.7)

# Checking performance
gbm_train_perf_orig = model_performance_classification_sklearn(tuned_gbm_orig, X_train_over, y_train_over)
print("Performance on Train Set:")
gbm_train_perf_orig

Performance on Train Set:

    Accuracy  Recall  Precision    F1
0      0.912   0.825      1.000 0.904

gbm_val_perf_orig = model_performance_classification_sklearn(tuned_gbm_orig, X_val, y_val)
print("Performance on Validation Set:")
gbm_val_perf_orig

Performance on Validation Set:

    Accuracy  Recall  Precision    F1
0      0.979   0.715      0.869 0.785

```

Hyperparameter Tuning of Gradient Boost on Oversampled Data

```

# Defining the Gradient Boosting model for oversampled data
gb_model_over = GradientBoostingClassifier(random_state=1)

# Parameter grid for oversampled data (you can adjust or use the same)
param_grid_over = {
    "n_estimators": np.arange(100, 150, 25),
    "learning_rate": [0.2, 0.05, 1],
    "subsample": [0.5, 0.7],
    "max_features": [0.5, 0.7]
}

# Setting up RandomizedSearchCV for the oversampled data
randomized_cv_over = RandomizedSearchCV(estimator=gb_model_over, param_distributions=param_grid_over,
                                       scoring=scorer, n_iter=50, n_jobs=-1, cv=5, random_state=1)

# Fitting RandomizedSearchCV on the oversampled data
randomized_cv_over.fit(X_train_over, y_train_over)

# Extracting best parameters for the oversampled data
best_params_over = randomized_cv_over.best_params_

# Creating a new Gradient Boosting Classifier with the best parameters for oversampled data
tuned_gbm_over = GradientBoostingClassifier(
    max_features=best_params_over['max_features'],

```

```

        random_state=1,
        learning_rate=best_params_over['learning_rate'],
        n_estimators=best_params_over['n_estimators'],
        subsample=best_params_over['subsample']
    )

    # Fitting the tuned model on the oversampled training data
    tuned_gbm_over.fit(X_train_over, y_train_over)

    GradientBoostingClassifier(learning_rate=1, max_features=0.5,
                               n_estimators=np.int64(125), random_state=1,
                               subsample=0.7)

    # Checking performance
    gbm_train_perf_over = model_performance_classification_sklearn(tuned_gbm_over, X_train_over, y_train_over)
    gbm_train_perf_over

    Accuracy  Recall  Precision    F1
0      0.993   0.992    0.993  0.993

    gbm_val_perf_over = model_performance_classification_sklearn(tuned_gbm_over, X_val, y_val)
    gbm_val_perf_over

    Accuracy  Recall  Precision    F1
0      0.968   0.867    0.657  0.748

```

Hyperparameter Tuning of Gradient Boost on Undersampled Data

```

# Defining the Gradient Boosting model for undersampled data
gb_model_un = GradientBoostingClassifier(random_state=1)

# Parameter grid for undersampled data (you can adjust or use the same)
param_grid_un = {
    "n_estimators": np.arange(100, 150, 25),
    "learning_rate": [0.2, 0.05, 1],
    "subsample": [0.5, 0.7],
    "max_features": [0.5, 0.7]
}

# Setting up RandomizedSearchCV for the undersampled data
randomized_cv_un = RandomizedSearchCV(estimator=gb_model_un, param_distributions=param_grid_un,
                                       scoring=scorer, n_iter=50, n_jobs=-1, cv=5, random_state=1)

# Fitting RandomizedSearchCV on the undersampled data
randomized_cv_un.fit(X_train_un, y_train_un)

# Extracting best parameters for the undersampled data
best_params_un = randomized_cv_un.best_params_

```

```

# Creating a new Gradient Boosting Classifier with the best parameters for undersampled data
tuned_gbm_un = GradientBoostingClassifier(
    max_features=best_params_un['max_features'],
    random_state=1,
    learning_rate=best_params_un['learning_rate'],
    n_estimators=best_params_un['n_estimators'],
    subsample=best_params_un['subsample']
)

# Fitting the tuned model on the undersampled training data
tuned_gbm_un.fit(X_train_un, y_train_un)

GradientBoostingClassifier(learning_rate=0.2, max_features=0.5,
                           n_estimators=np.int64(100), random_state=1,
                           subsample=0.5)

# Checking performance
gbm_train_perf_un = model_performance_classification_sklearn(tuned_gbm_un, X_train_un, y_train_un)
gbm_train_perf_un

    Accuracy  Recall  Precision    F1
0      0.984   0.975     0.993 0.984

gbm_val_perf_un = model_performance_classification_sklearn(tuned_gbm_un, X_val, y_val)
gbm_val_perf_un

    Accuracy  Recall  Precision    F1
0      0.909   0.878     0.361 0.511

```

Hyperparameter Tuning of AdaBoost on Original Data

```

# Defining the AdaBoost model for original data
ada_model_orig = AdaBoostClassifier(random_state=1)

# Parameter grid for AdaBoost
param_grid_ada_orig = {
    "n_estimators": [50, 100, 150],
    "learning_rate": [0.01, 0.05],
}

# Setting up RandomizedSearchCV for the original data
randomized_ada_orig = RandomizedSearchCV(estimator=ada_model_orig, param_distributions=param_grid_ada_orig,
                                         scoring= scorer, n_iter=10, cv=5, random_state=1)

# Fitting RandomizedSearchCV on the original data
randomized_ada_orig.fit(X_train, y_train)

```



```

# Extracting best parameters
best_params_ada_orig = randomized_ada_orig.best_params_

# Creating a new pipeline of AdaBoost Classifier with the best parameters
tuned_ada_orig = AdaBoostClassifier(**best_params_ada_orig, random_state=1)

# Fitting the tuned model on the original training data
tuned_ada_orig.fit(X_train, y_train)

AdaBoostClassifier(learning_rate=0.05, n_estimators=150, random_state=1)

# Evaluating the tuned model
ada_train_perf_orig = model_performance_classification_sklearn(tuned_ada_orig, X_train, y_train)
ada_train_perf_orig

    Accuracy  Recall  Precision    F1
0      0.951   0.140     0.915 0.244

ada_val_perf_orig = model_performance_classification_sklearn(tuned_ada_orig, X_val, y_val)
ada_val_perf_orig

    Accuracy  Recall  Precision    F1
0      0.952   0.137     0.881 0.237

```

Hyperparameter Tuning of AdaBoost on Oversample Data

```

# Defining the AdaBoost model for oversampled data
ada_model_over = AdaBoostClassifier(random_state=1)

# Parameter grid for AdaBoost
param_grid_ada_over = {
    "n_estimators": [50, 100, 150],
    "learning_rate": [0.01, 0.05],
}

# Setting up RandomizedSearchCV for the oversampled data
randomized_ada_over = RandomizedSearchCV(estimator=ada_model_orig, param_distributions=param_grid_ada_over,
                                         scoring=scorer, n_iter=10, cv=5, random_state=1)

# Fitting RandomizedSearchCV on the oversampled data
randomized_ada_over.fit(X_train_over, y_train_over)

# Extracting best parameters for the oversampled data
best_params_ada_over = randomized_ada_over.best_params_

# Creating a new AdaBoost Classifier with the best parameters for oversampled data
tuned_ada_over = AdaBoostClassifier(**best_params_ada_over, random_state=1)

```

```

# Fitting the tuned model on the oversampled training data
tuned_ada_over.fit(X_train_over, y_train_over)

AdaBoostClassifier(learning_rate=0.01, random_state=1)

print("Best parameters are {} with CV score={}" .format(randomized_ada_over.best_params_, randomized_ada_over.best_cv_score_))
Best parameters are {'n_estimators': 50, 'learning_rate': 0.01} with CV score=0.851694915254

# Evaluating the tuned model on oversampled data
ada_train_perf_over = model_performance_classification_sklearn(tuned_ada_over, X_train_over, y_train_over)
ada_train_perf_over

    Accuracy  Recall  Precision    F1
0      0.759   0.864     0.714 0.782

ada_val_perf_over = model_performance_classification_sklearn(tuned_ada_over, X_val, y_val)
ada_val_perf_over

    Accuracy  Recall  Precision    F1
0      0.664   0.804     0.118 0.205

```

Hyperparameter Tuning of AdaBoost on Undersampled Data

```

# Defining the AdaBoost model for oversampled data
ada_model_un = AdaBoostClassifier(random_state=1)

# Parameter grid for AdaBoost
param_grid_ada_un = {
    "n_estimators": [50, 100, 150],
    "learning_rate": [0.01, 0.05],
}

# Setting up RandomizedSearchCV for the undersampled data
randomized_ada_un = RandomizedSearchCV(estimator=ada_model_orig, param_distributions=param_grid_ada_un,
                                       scoring=scorer, n_iter=10, cv=5, random_state=1)

# Fitting RandomizedSearchCV on the undersampled data
randomized_ada_un.fit(X_train_un, y_train_un)

# Extracting best parameters for the undersampled data
best_params_ada_un = randomized_ada_un.best_params_

# Creating a new AdaBoost Classifier with the best parameters for undersampled data
tuned_ada_un = AdaBoostClassifier(**best_params_ada_un, random_state=1)

# Fitting the tuned model on the undersampled training data
tuned_ada_un.fit(X_train_un, y_train_un)

AdaBoostClassifier(learning_rate=0.05, n_estimators=150, random_state=1)

```

```

# Evaluating the tuned model on undersampled data
ada_train_perf_un = model_performance_classification_sklern(tuned_ada_un, X_train_un, y_train_un)
ada_train_perf_un

    Accuracy  Recall  Precision    F1
0      0.854   0.781      0.915 0.843

ada_val_perf_un = model_performance_classification_sklern(tuned_ada_un, X_val, y_val)
ada_val_perf_un

    Accuracy  Recall  Precision    F1
0      0.898   0.733      0.312 0.438

```

Model Performance Comparison

```

# Training performance comparison

```

```

models_train_comp_df = pd.concat(
    [
        rf_train_perf_orig.T,
        gbm_train_perf_orig.T,
        ada_train_perf_orig.T,

        rf_train_perf_over.T,
        gbm_train_perf_over.T,
        ada_train_perf_over.T,

        rf_train_perf_un.T,
        gbm_train_perf_un.T,
        ada_train_perf_un.T,

    ],
    axis=1,
)
models_train_comp_df.columns = [
    'rf_train_perf_orig',
    'gbm_train_perf_orig',
    'ada_train_perf_orig',

    'rf_train_perf_over',
    'gbm_train_perf_over',
    'ada_train_perf_over',

    'rf_train_perf_un',
    'gbm_train_perf_un',
    'ada_train_perf_un',
]

```

```

]
print("Training performance comparison:")
models_train_comp_df

Training performance comparison:

      rf_train_perf_orig  gbm_train_perf_orig  ada_train_perf_orig  \
Accuracy                0.993                0.912                0.951
Recall                  0.879                0.825                0.140
Precision                0.999                1.000                0.915
F1                      0.935                0.904                0.244

      rf_train_perf_over  gbm_train_perf_over  ada_train_perf_over  \
Accuracy                0.999                0.993                0.759
Recall                  0.998                0.992                0.864
Precision                1.000                0.993                0.714
F1                      0.999                0.993                0.782

      rf_train_perf_un  gbm_train_perf_un  ada_train_perf_un
Accuracy              0.966              0.984              0.854
Recall                0.940              0.975              0.781
Precision              0.991              0.993              0.915
F1                    0.965              0.984              0.843

# Validation performance comparison

models_val_comp_df = pd.concat(
    [
        rf_val_perf_orig.T,
        gbm_val_perf_orig.T,
        ada_val_perf_orig.T,

        rf_val_perf_over.T,
        gbm_val_perf_over.T,
        ada_val_perf_over.T,

        rf_val_perf_un.T,
        gbm_val_perf_un.T,
        ada_val_perf_un.T,
    ],
    axis=1,
)

models_val_comp_df.columns = [
    'rf_val_perf_orig',
    'gbm_val_perf_orig',
    'ada_val_perf_orig',

```

```

        'rf_val_perf_over',
        'gbm_val_perf_over',
        'ada_val_perf_over',

        'rf_val_perf_un',
        'gbm_val_perf_un',
        'ada_val_perf_un',
    ]
    print("Validation performance comparison:")
    models_val_comp_df

```

Validation performance comparison:

	rf_val_perf_orig	gbm_val_perf_orig	ada_val_perf_orig	\
Accuracy	0.982	0.979	0.952	
Recall	0.670	0.715	0.137	
Precision	0.989	0.869	0.881	
F1	0.799	0.785	0.237	

	rf_val_perf_over	gbm_val_perf_over	ada_val_perf_over	\
Accuracy	0.988	0.968	0.664	
Recall	0.859	0.867	0.804	
Precision	0.917	0.657	0.118	
F1	0.887	0.748	0.205	

	rf_val_perf_un	gbm_val_perf_un	ada_val_perf_un	
Accuracy	0.935	0.909	0.898	
Recall	0.878	0.878	0.733	
Precision	0.450	0.361	0.312	
F1	0.595	0.511	0.438	

```

    print("Training performance comparison:")
    models_train_comp_df

```

Training performance comparison:

	rf_train_perf_orig	gbm_train_perf_orig	ada_train_perf_orig	\
Accuracy	0.993	0.912	0.951	
Recall	0.879	0.825	0.140	
Precision	0.999	1.000	0.915	
F1	0.935	0.904	0.244	

	rf_train_perf_over	gbm_train_perf_over	ada_train_perf_over	\
Accuracy	0.999	0.993	0.759	
Recall	0.998	0.992	0.864	
Precision	1.000	0.993	0.714	
F1	0.999	0.993	0.782	

	rf_train_perf_un	gbm_train_perf_un	ada_train_perf_un
Accuracy	0.966	0.984	0.854
Recall	0.940	0.975	0.781
Precision	0.991	0.993	0.915
F1	0.965	0.984	0.843

```
print("Validation performance comparison:")
models_val_comp_df
```

Validation performance comparison:

	rf_val_perf_orig	gbm_val_perf_orig	ada_val_perf_orig	\
Accuracy	0.982	0.979	0.952	
Recall	0.670	0.715	0.137	
Precision	0.989	0.869	0.881	
F1	0.799	0.785	0.237	

	rf_val_perf_over	gbm_val_perf_over	ada_val_perf_over	\
Accuracy	0.988	0.968	0.664	
Recall	0.859	0.867	0.804	
Precision	0.917	0.657	0.118	
F1	0.887	0.748	0.205	

	rf_val_perf_un	gbm_val_perf_un	ada_val_perf_un
Accuracy	0.935	0.909	0.898
Recall	0.878	0.878	0.733
Precision	0.450	0.361	0.312
F1	0.595	0.511	0.438

AdaBoost on oversampled data emerges as a strong candidate for the final model due to its excellent performance across accuracy, **recall**, precision, and F1 score in the validation set.

Final Model Selection

Check the performance of the final model on the test data


```
ada_test = model_performance_classification_sklearn(tuned_ada_over, X_test, y_test)
ada_test
```

	Accuracy	Recall	Precision	F1
0	0.675	0.840	0.131	0.226

```
feature_names = X_train.columns
importances = tuned_ada_over.feature_importances_
indices = np.argsort(importances)
```

```
plt.figure(figsize=(10, 10))
plt.title("Feature Importances")
```

```
plt.barh(range(len(indices)), importances[indices], color="violet", align="center")
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel("Relative Importance")
plt.show()
```



Pipelines to build the final model

Since we have only one datatype in the data we don't need to use column transformer which can be used to personalize the pipeline to perform different preprocessing steps on different columns. Note: see case study how to use this.

The steps followed in implementing a pipeline are -

- Pipeline() - A pipeline is defined as a list of tuples
- Pipeline.fit() - Pipeline is fitted on the train set
- Pipeline.score() - Pipeline objects checks the performance.

```

pipe = Pipeline(
    steps=[
        ("imputer", SimpleImputer(strategy="median")),
        ("scaler", StandardScaler()),
        ("AdaBoost", AdaBoostClassifier(
            (
                n_estimators= 50,
                learning_rate= 0.05,
            ),
        )),
    ]
)

# Check pipeline steps
pipe.steps

[('imputer', SimpleImputer(strategy='median')),
 ('scaler', StandardScaler()),
 ('AdaBoost', AdaBoostClassifier(learning_rate=0.05))]

# Fit the pipeline on training data as if it were a model since the model is included in it
pipe.fit(X_train, y_train)
# See pipeline structure below

Pipeline(steps=[('imputer', SimpleImputer(strategy='median')),
                 ('scaler', StandardScaler()),
                 ('AdaBoost', AdaBoostClassifier(learning_rate=0.05))])

```

The above means a model is built which becomes part of pipe

```

# pipe object's accuracy on the train set
pipe.score(X_train, y_train)

```

0.944

Instead of model.score calling pipe.score function does a predict then uses those results with the actual results to give the score.

```

# pipe object's accuracy on the test set
pipe.score(X_test, y_test)

```

0.9436

The data called by pipe.score() undergoes transformation process in the first two steps (imputer and scaler) and the predict function in the 3rd step ("AdaBoost").

Questions I still have:

- Am I suppose to fit the pipeline into the X_train, y_train or X and y?
- What about X_train_over, y_train_over?
- Is SMOTE incorporated into the pipeline?

Business Insights and Conclusions

Features V18, V39, and V12 have significant influence on the model's predictions.

The company can focus on further analyzing these important features to better understand their underlying relationships with the target variable.